

Overnight digest

Friday 14 August 2026, 15:45 · 9 session(s) · 8 with blockers · 9 awaiting a decision

AgentMemoryBench d533133b

PROGRESS

- Published PersonaMem + PrefEval to S3, Athena, and the HTML dashboard (commit 6eb2395d); Contexto ties AgentCore on PrefEval (0.593 vs 0.667, $p=0.108$), loses on PersonaMem (0.324 vs 0.491).
- Built a matched 4-arm trace-comparison view (artifacts/personalization/trace_compare.html) showing WHY they differ: AgentCore stores structured preference objects (4-5 retrieved), Contexto stores atomic facts (1-2 retrieved).
- Closed the LoCoMo AgentCore arm (0.367, beats keepstate 0.331) and the full_history-1M run.

BLOCKERS

none. All four tracked tasks closed; nothing running on the box.

DECISIONS FOR FABIO

none. The one open call — re-run full_history-1M for \$1.5K to get a clean-store number — Fabio already decided (option 2: no re-run; the collapse finding stands on the 200K campaign).

Strands 17cbae8f

PROGRESS

- Wired the self-hosted model path (``vllm/<model>`` → Strands ``OpenAIModel`` passed prebuilt into Stan), launched a spot g6e.12xlarge (4xL40S), and served Qwen3-8B on vLLM after patching a flashinfer py3.10 bug and a missing ``ninja``.
- Ran the untrained Qwen3-8B eval (140x2): baseline 0.7%, BPE 0.0% — the pre-training "before" row.
- Built and validated the replay-based SFT dataset: 4,458 (observation + BPE-state → action) pairs from opus5 (90.0% recovery) + sonnet5 (86.7%); 12 Bedrock reference episodes confirm replay is lossless (100% observation similarity, 162/162 steps exact). Logged all of it to an MLflow dashboard (live on port 5001).

BLOCKERS

none — all four overnight jobs (replay x2, reference, untrained eval) finished and results are on S3.

DECISIONS FOR FABIO

Whether to start LoRA-SFT (task #4) now on the 4,458 pairs, or wait for Fabio to inspect the MLflow dashboard and sign off on dataset quality (including the ~10-13% unrecovered episodes) first.

swe-bench-track 6326dc1b

PROGRESS

- Scored all 500 SWE-bench Verified instances through the Intent Gap pipeline: 185 low / 275 medium / 25 high (median 0.30, max 0.72). The mining track produced 0 high-gap from 56 tasks; this is the result that justifies

Contexto 0baea953

PROGRESS

- Pressure-tested spec P1 on 149 real sessions: 0/149 overflow. Found the real defect — the text-only filter drops 71% of content (all tool calls). Wrote docs and a spec appendix correcting P1 to content-type

the second track.

- Stood up the Seam-3 grading box (EC2 `i-022538b05b2cdc9fb`, stack `SweGradingStack`, isolated from the running benchmark box) with SWE-bench harness 4.1.0; gold-patch check passed 2/2, confirming the box reproduces SWE-bench pass/fail.
- Cherry-picked the trace-explorer `</script>`-injection fix to main; committed the full second-track work (adapter, scorer, grader, calibration tooling) to branch `worktree-swe-bench-track`.

BLOCKERS

none. (500-instance adapt ran 0 skips; scoring completed all 12 repos; grading box online.)

DECISIONS FOR FABIO

- Merge `worktree-swe-bench-track` into main. Held back deliberately: main is being actively committed to by the running-benchmark session, and the one overlapping file (`build\_trace\_app.py`) is a clean superset, so the merge is trivial but is a shared-line call to make.
- Whether to cut the SWE IG calibration sample for you to label. The 500 IG scores are the pipeline's, uncalibrated on SWE-bench data; the generator is built and tested, waiting on your go.

warpsessions fcc26a14

PROGRESS

- Built `bus.py`: a file-based inter-session bus letting live Claude sessions ask each other questions (1:1 or broadcast); a daemon answers by forking the target session (`claude -p --resume --fork-session`) so the live session is never interrupted. Added the `bus` skill and verified end-to-end (Contexto and AgentMemoryBench forks answered from real context).
- Built `manager.py`: a local web console

filtering.

- Built the P2 context-field prototype on branch `p2-context-prototype` (pushed, 569 tests green); the AgentMemoryBench session ran the PersonaMem gate with it.
- Fixed a real production bug the gate surfaced: `\_MAX\_TOKENS=1024` truncated extractor JSON on fact-dense conversations, silently dropping the whole batch. Raised to 4096 + fail-loud on truncation. Branch `fix-extract-output-truncation`, pushed, 571 tests green.

BLOCKERS

- P2 PersonaMem gate is partial: canon and context arms done (0.3565 vs 0.3599, +0.34pp), agentcore arm still running (slow by design, ~60s/conversation — not stuck). Final McNemar verdict waits on it.

DECISIONS FOR FABIO

- Merge `fix-extract-output-truncation` to mainline? It is a clean standalone bug fix, tests green, waiting on your go (or your self-merge stance applies).
- P2 build/drop decision comes once agentcore lands; early read (+0.34pp vs a 50% pre-registered bar) trends toward drop.

MemorySkillServiceDocs-ws dbca4bbe

PROGRESS

- Ran GG3b (forced skill consultation on the 12 frozen tau2 tasks): engagement fired for the first time (7/12 reads), and the skill body demonstrably moved behavior both ways — task 33 fixed, task 111 regressed; committed with corrected numbers.
- Ran GG4, the 40-task retail ruler (2 arms × 3 repeats): control 0.750 stable, disclosure-forced 0.692; 39/40 tasks deterministic, so future gates drop to ~41 sims; skills are

(localhost:8799) to chat with a persistent manager session that lists live sessions, browses their Claude memory files, and runs the bus. Added `bus.py inject` (types a prompt into a live session's Warp tab, with a dead-pane guard) as the manager's write path.

- Hardened the bus for stale session ids: reroute a 1:1 question to the live session in the same project dir, else write an UNDELIVERABLE NACK after 60s. Both paths tested.

- Reworked the menu bar app's state marks (● blocked, * working, no mark = waiting) via a Notification hook writing per-session state files.

- Published the repo: ``github.com/nonatofabio/warpsessions``, sanitized of personal references, flipped public.

BLOCKERS

none. (Possible transient Bedrock 503s were noted while testing; this fork answered, so the API is reachable.)

DECISIONS FOR FABIO

`slack-note.md` remains in git history after being untracked — Fabio to decide whether to rewrite history with `git filter-repo` and force-push, or leave it (judged harmless).

strongly bidirectional (4 clean wins, 5 clean losses); the co-request-drop defect now reproduces 6/6 on tasks 111+42, the top mutation target.

- Wrote SKILL-CREATION-OPTIMIZATION.md (mine → gate → optimize loop), incorporated all 10 of Fabio's Mindle comments, and stood up the results dashboard by reusing another session's MLflow component.

BLOCKERS

- MLflow store on the devbox is down (alembic migration error, `Can't locate revision 6f8d9c3b2a1e`); another agent owns the fix. GG4 results are committed in git and will backfill in under a minute once the store is back — not on the critical path.

DECISIONS FOR FABIO

- none waiting. Fabio already approved the queue (40-task ruler done; next: loss forensics on tasks 64/68/97, then the task-111/42 mutation cycle, then the SkillOpt spike). Second gate arm settled as a coding task (BigCodeBench candidate).

tetra a6f71ace

PROGRESS

- Delivered three documents to ``/Users/fnp/Documents/wd/dev/tetra`` proposing a learned memory policy loop for Tetra Build: a one-pager, an AWSTetra-style SPEC (``docs/specs/2026/08/memory-policy-loop``), and a scientific foundations doc. All three synthesize four sources (ACM 2607.23809, agent-compatible compression 2605.30785, proactive memory agent 2607.08716, rLLM realtime RL) plus Strands and the Velocity strategy docs.

cantjump 140cfe0a

PROGRESS

- Rebutted "LLMs Can't Jump" with a live 70-cell experiment: 14 models, 7 families, invented laws (three-fold gravity, ``popcount(x⊕(x>>3))``) so nothing could be recalled. Opus 4.8, Sonnet 5, DeepSeek V3.2, and GPT-OSS-120B made the jump (holdout errors $\sim 10^{-5}$; oracle 20/20).
- Added the three controls that turn the demo into an argument: null universe (catches confabulation), passive/no-tools ablation (proves interaction is causally necessary), and

- Ran the humanize and amazon-writing skills on all three: sentence-length CV 0.53–0.84, zero em-dashes, bold-colon lists converted to prose; opened all three in Mindle for review.
- Core framing landed: Tetra's event ledger + evidence manifest are the trajectory record and reward signal an RL memory loop needs; phase-0 replay experiment is scoped to two people / two weeks / zero training compute.

BLOCKERS

none

DECISIONS FOR FABIO

- Three items await your review, none blocking: (1) read-aloud/voice pass on the docs was not done, flagged for you to do in Mindle; (2) reward-hacking mitigations in the SPEC are listed as open (holdout critics, mutation-weighted gates); (3) where the memory policy runs — Strands-hosted vs inside `agent-runtime` — is left open pending your call.

a size gradient across every family. All I1–I5 integrity checks pass over 1,360 readings + 517 oracle queries.

- Wrote `rebuttal\_v2.md` via the humanize skill (v1 preserved); audit clean, cadence CV 0.81, voice-matched.

BLOCKERS

none. All 70 cells verified; the one zero-evidence null pass (GPT-OSS-20B) is already flagged ✓* and excluded.

DECISIONS FOR FABIO

none blocking. Optional when you're back: whether to publish v1 or v2 as the canonical paper, and whether to add more seeds per cell (currently 2 on `main`, 1 elsewhere — enough for the existential claim, not for success rates).

intent-benchmark 874b5c41

PROGRESS

- Completed the full two-repo cap-128 grid (pytest+cargo, 62 reportable triples: i 40 / ii 46 / iii 47), seed-replicated all 6 pytest discordants (iii 12/12 across seeds), read the two anti-iii cargo flips (one was an oracle flake, removed; one a non-querying miss), and wrote the Aug-14 milestone numbers doc.
- Shipped compiler v3 (process-claim filter, quote cleanup, pointer summaries), archived the compiler-v2 grid as the reference baseline, and launched the multi-repo grid over the 183-task pool on two lanes.
- Multi-repo grid at ~117 rows: django 75 rows / 23 triples with 4 monotone discordants (2 iii-only rescues), vinext 49 rows — the iii-rescue pattern now replicates across pytest, django, and vinext (3 languages).

BLOCKERS

none. Midway on the box lapsed twice overnight but cached tokens bridged both gaps with zero lost runs; refreshed now. Disk hit 82%, reclaimed to 78% by pruning recorded run-trees.

DECISIONS FOR FABIO

none pending. (aider and ruff lanes queued behind vinext/django; full six-repo consolidated report auto-runs when the last repo completes.)

Regenerated every run. Data: overnight/history/*.json · Sessions forked read-only — live sessions untouched.